

Final Project

```
library(imputeTS)

## Warning: package 'imputeTS' was built under R version 3.3.3

library(forecast)

## Warning: package 'forecast' was built under R version 3.3.3

library(zoo)

## Warning: package 'zoo' was built under R version 3.3.3

##
## Attaching package: 'zoo'

## The following object is masked from 'package:imputeTS':
##
##     na.locf

## The following objects are masked from 'package:base':
##
##     as.Date, as.Date.numeric

library(ggplot2)

## Warning: package 'ggplot2' was built under R version 3.3.3
```

Read in CO2 data

```
co2 <- read.csv("monthly co2.csv")
```

Explore data

```
str(co2)

## 'data.frame':    713 obs. of  7 variables:
## $ Year          : int  1958 1958 1958 1958 1958 1958 1958 1958 1958 1958
## $ Month         : int   3  4  5  6  7  8  9 10 11 12 ...
## $ decimal.date   : num  1958 1958 1958 1958 1959 ...
## $ average        : num  316 317 318 -100 316 ...
## $ interpolated   : num  316 317 318 317 316 ...
## $ trend..season.corr.: num  315 315 315 315 315 ...
## $ X.days         : int   -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 ...

head(co2)

##   Year Month decimal.date average interpolated trend..season.corr. X.days
## 1 1958     3    1958.208   315.71         315.71         314.62      -1
## 2 1958     4    1958.292   317.45         317.45         315.29      -1
## 3 1958     5    1958.375   317.50         317.50         314.71      -1
```

##	4	1958	6	1958.458	-99.99	317.10	314.85	-1
##	5	1958	7	1958.542	315.86	315.86	314.98	-1
##	6	1958	8	1958.625	314.93	314.93	315.94	-1

Trim to CO2 column and data since 1966

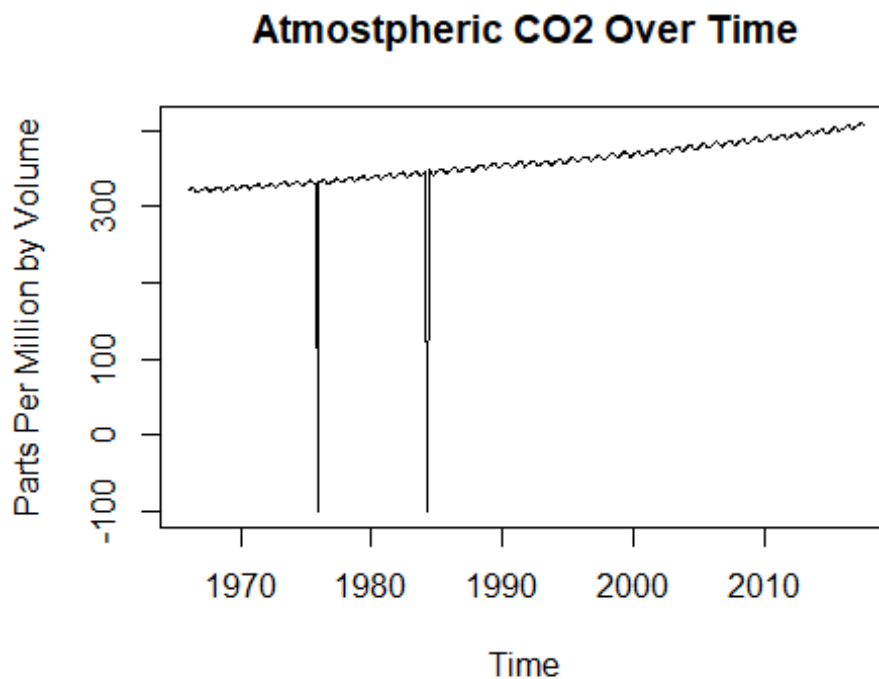
```
co2 <- co2[,c("Year", "Month", "average")]
colnames(co2) <- c("Year", "Month", "CO2")
co2 <- co2[co2$Year >= 1966, ]
co2.ts <- ts(co2$CO2, start = 1966, frequency = 12)
head(co2.ts)
```

```
##           Jan      Feb      Mar      Apr      May      Jun
## 1966 320.62 321.59 322.39 323.87 324.01 323.75
```

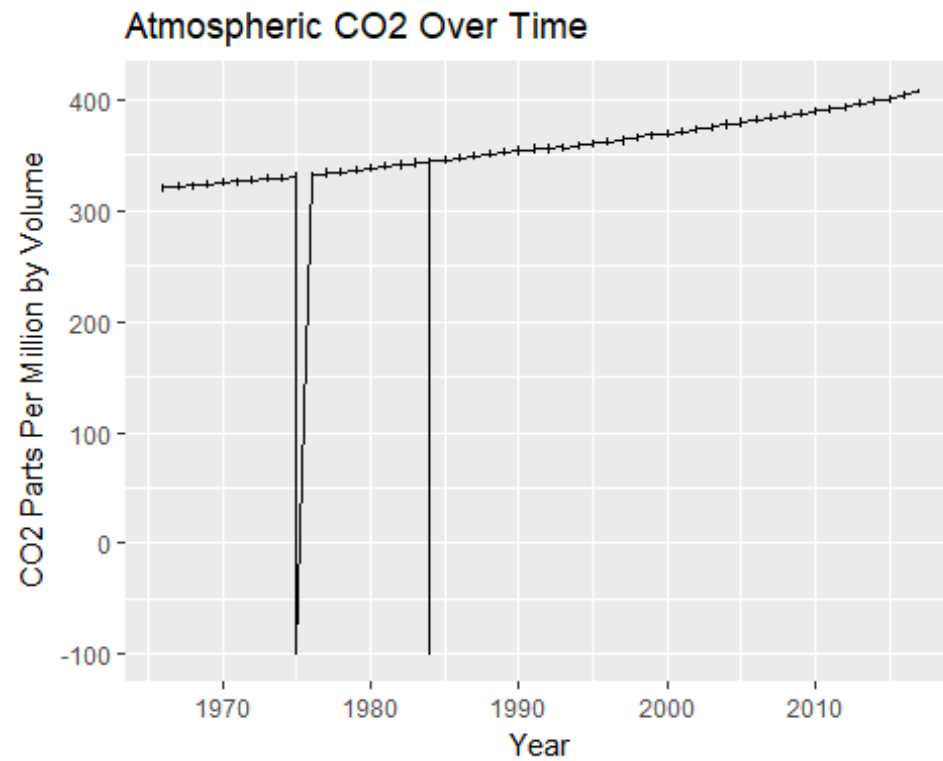
Plot time series

Observe any missing values or outliers

```
#plotNA.distribution(co2.ts)
plot.ts(co2.ts, main = "Atmospheric CO2 Over Time", ylab = "Parts Per
Million by Volume")
```



```
ggplot(co2, aes(x=Year, y=CO2)) +
  geom_line() +
  labs(y = "CO2 Parts Per Million by Volume", title = "Atmospheric CO2 Over
Time")
```




Data < 0 for atmospheric CO2 doesn't make sense so we will remove these and impute data.

```
co2.ts <- ifelse(co2.ts < 0, NA, co2.ts)
```

```
plotNA.distribution(co2.ts)
```


Distribution of NAs



Impute missing values

```
co2.ts.imp <- na.interpolation(co2.ts)
```

```
plotNA.distribution(co2.ts.imp)
```





Conclusions: Plot of CO2 time series indicates positive, additive trend and annual seasonality.

ACF

```
#acf(co2.ts.imp, lag = 100)
# figure
# subplot(2,1,1)
# autocorr(co2.ts.imp)
# subplot(2,1,2)
# parcorr(co2.ts.imp)


par(mfrow=c(2,1))
plot(acf(co2.ts.imp, plot=FALSE), ylab = "CO2 Autocorrelation", main = "")
pacf(co2.ts.imp)
```



Long memory auto correlation. The persistence of high values in acf plot indicate a long term positive trend. These results are consistent with our plot of the time series.

PACF

```
pacf(co2.ts.imp, main = "", ylab = "CO2 Partial Autocorrelation")
```



Long memory

partial autocorrelation drops off around lag 20.

why doesn't x adjust with lag?

Training and test

```
co2.ts.train <- ts(co2.ts.imp[1:588], start = 1966, frequency = 12) # through 2014
co2.ts.test <- ts(co2.ts.imp[589:600], start = 2015, frequency = 12) # all of 2015
```

Seasonal ARIMA

Seasonal time series

```
sarima.mod <- auto.arima(co2.ts.train) # model through 2014
summary(sarima.mod)

## Series: co2.ts.train
## ARIMA(0,1,2)(1,1,2)[12]
##
## Coefficients:
##
## Warning in sqrt(diag(x$var.coef)): NaNs produced
##          ma1      ma2      sar1      sma1      sma2
##      -0.3543 -0.0639 -0.6315 -0.2541 -0.5608
## s.e.   0.0421  0.0423      NaN      NaN      NaN
##
```


```
## sigma^2 estimated as 0.09258: log likelihood=-131.22
## AIC=274.44 AICc=274.58 BIC=300.56
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE
## Training set 0.0220546 0.2995792 0.2327383 0.005981288 0.06555839
##              MASE      ACF1
## Training set 0.1443602 0.008156477
```

Seasonal ARIMA Forecast

```
(sarima.2015 <- forecast(sarima.mod, h=12))
```

```
##      Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## Jan 2015      400.0454 399.6555 400.4354 399.4491 400.6418
## Feb 2015      400.8055 400.3413 401.2696 400.0956 401.5153
## Mar 2015      401.7118 401.1951 402.2284 400.9216 402.5019
## Apr 2015      403.0050 402.4407 403.5693 402.1420 403.8680
## May 2015      403.6312 403.0231 404.2394 402.7011 404.5614
## Jun 2015      402.9493 402.3002 403.5984 401.9566 403.9421
## Jul 2015      401.3714 400.6837 402.0590 400.3197 402.4230
## Aug 2015      399.2930 398.5689 400.0171 398.1856 400.4004
## Sep 2015      397.8277 397.0689 398.5865 396.6672 398.9881
## Oct 2015      397.9847 397.1927 398.7767 396.7735 399.1959
## Nov 2015      399.3993 398.5755 400.2232 398.1394 400.6593
## Dec 2015      400.8738 400.0193 401.7283 399.5669 402.1807
```

```
plot(co2.ts.imp[1:600])
lines(sarima.2015$mean, col = "green")
```

Calculate error...

Holt-Winters

Time series with positive trend and seasonality

```
hw.co2.mod <- HoltWinters(co2.ts.train)
summary(hw.co2.mod)
```

##		Length	Class	Mode
## fitted		2304	mts	numeric
## x		588	ts	numeric
## alpha		1	-none-	numeric
## beta		1	-none-	numeric
## gamma		1	-none-	numeric
## coefficients		14	-none-	numeric
## seasonal		1	-none-	character
## SSE		1	-none-	numeric
## call		2	-none-	call


```
hw.co2.mod

## Holt-Winters exponential smoothing with trend and additive seasonal
component.
##
## Call:
## HoltWinters(x = co2.ts.train)
##
```

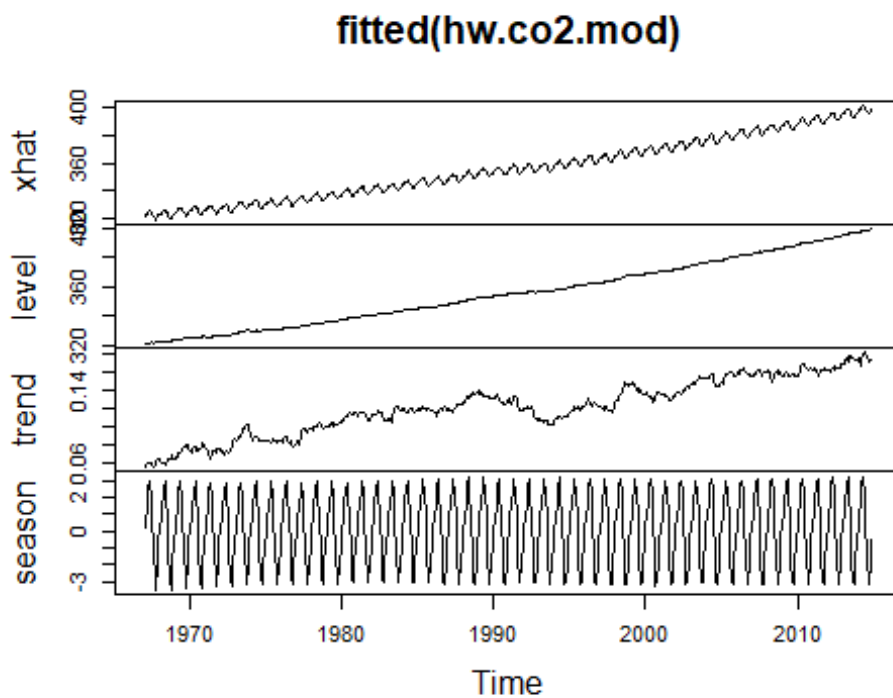
```
## Smoothing parameters:
## alpha: 0.5460384
## beta : 0.01522105
## gamma: 0.2959677
##
## Coefficients:
##          [,1]
## a  399.4284106
## b    0.1751974
## s1   0.4492880
## s2   1.0135282
## s3   1.7472395
## s4   2.8265649
## s5   3.2747156
## s6   2.3369992
## s7   0.6081790
## s8  -1.5913489
## s9  -3.1618090
## s10 -3.1275772
## s11 -1.8776697
## s12 -0.5626363
```

```
plot(hw.co2.mod)
```

Holt-Winters filtering



```
plot(fitted(hw.co2.mod))
```



```
# forecast
```

Read in annual temps

```
temps <- read.csv("GlobalTemperatures.csv")
#temps.ho <- read.csv("16-17_temps.csv")

names(temps)

## [1] "dt"
## [2] "LandAverageTemperature"
## [3] "LandAverageTemperatureUncertainty"
## [4] "LandMaxTemperature"
## [5] "LandMaxTemperatureUncertainty"
## [6] "LandMinTemperature"
## [7] "LandMinTemperatureUncertainty"
## [8] "LandAndOceanAverageTemperature"
## [9] "LandAndOceanAverageTemperatureUncertainty"

temps <- temps[,c("dt", "LandAndOceanAverageTemperature")]
temps$dt <- as.character(temps$dt)
```

Grab only data since 1966

```
temps <- temps[temps$dt >= "1966-01-01",]
colnames(temps) <- c("dt", "temp")
```

Streamline holdout data - 2015

```
# temps.ho$dt <- as.Date(as.yearmon(paste(temps.ho$year, temps.ho$month,
"01", sep = "-"))) # get date format
# temps.ho$dt <- as.character(temps.ho$dt) # match type in temps df
# temps.ho <- temps.ho[,c("dt", "raw.temp")]
# colnames(temps.ho) <- c("dt", "temp")
# temps.full <- rbind(temps, temps.ho) # make a full list
# temps.full.ts <- ts(temps.full$temp, start = 1966, frequency = 12)


temps.ts <- ts(temps$temp, start = 1966, frequency = 12)
temps.train <- temps[temps$dt <= "2014-12-01",]
temps.test <- temps[temps$dt >= "2015-01-01",]
```

Make Training Time Series

```
temps.train.ts <- ts(temps.train$temp, start = 1966, frequency = 12)
```


Plot missing values

```
plotNA.distribution(temps.ts, main = "Global Temperature Over Time", ylab =
"Temperature (°C)")
```




```
acf(temps$temp, main = "Temperature Autocorrelation")
```

Temperature Autocorrelation



```
pacf(temps$temp, main = "Temperature Partial Autocorrelation", lag = 50) #  
these are better
```

Temperature Partial Autocorrelation



Plot both TS together

```
# library(grid)
# library(dplyr)
#
# co2$dt <- as.Date(as.yearmon(paste(co2$Year, co2$Month, "01", sep = "-")))
# get dt column in co2
# plot1 <- co2 %>%
#   select(dt, CO2) %>%
#   na.omit() %>%
#   ggplot() + geom_line(aes(x = dt, y = CO2)) + ylab("Atmospheric CO2 ppm")
#   +
#   theme_minimal() +
#   theme(axis.title.x = element_blank())
#
# plot2 <- temps %>%
#   select(dt, temp) %>%
#   ggplot() + geom_line(aes(x = dt, y = temp)) + ylab("Temperature (°C)") +
#   theme_minimal() +
#   theme(axis.title.x = element_blank())
#
# grid.newpage()
# grid.draw(rbind(ggplotGrob(plot1), ggplotGrob(plot2), size = "last"))

##qplot(temps.ts) + geom_line(y=temps.ts)

#ggplot(temps, aes(x=dt, y=temp)) + geom_line()
```

Seasonal ARIMA

Seasonal time series

```
sarima.temps.mod <- auto.arima(temps.train.ts) # model through 2015
summary(sarima.temps.mod)

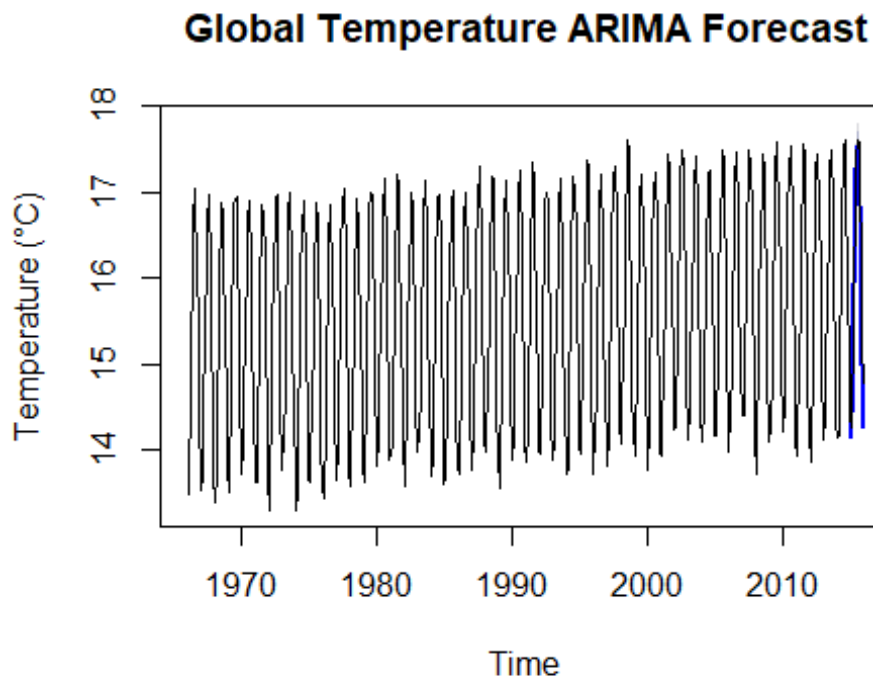
## Series: temps.train.ts
## ARIMA(2,0,2)(1,1,1)[12]
##
## Coefficients:
##          ar1      ar2      ma1      ma2      sar1      sma1
##      0.0556  0.7969  0.4013 -0.3348 -0.1909 -0.8239
## s.e.  0.0883  0.0862  0.1007  0.0731  0.0485  0.0352
##
## sigma^2 estimated as 0.009847:  log likelihood=507.76
## AIC=-1001.51   AICc=-1001.31   BIC=-971.02
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE
## Training set 0.01093378 0.09770333 0.07721149 0.0664059 0.5061589
##              MASE      ACF1
## Training set 0.5345283 0.008692596
```

Seasonal ARIMA Forecast


```
(sarima.temps.2015 <- forecast(sarima.temps.mod, h=12))
```

##		Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	Jan 2015	14.15179	14.02462	14.27897	13.95730	14.34629
##	Feb 2015	14.42931	14.28949	14.56913	14.21548	14.64314
##	Mar 2015	15.06586	14.91291	15.21881	14.83194	15.29977
##	Apr 2015	15.93116	15.77033	16.09200	15.68519	16.17714
##	May 2015	16.66667	16.49758	16.83576	16.40807	16.92527
##	Jun 2015	17.28573	17.11137	17.46009	17.01907	17.55239
##	Jul 2015	17.54924	17.36942	17.72905	17.27423	17.82424
##	Aug 2015	17.46753	17.28408	17.65098	17.18696	17.74810
##	Sep 2015	16.86487	16.67771	17.05202	16.57864	17.15110
##	Oct 2015	15.97062	15.78091	16.16033	15.68048	16.26076
##	Nov 2015	14.99596	14.80369	15.18822	14.70191	15.29000
##	Dec 2015	14.26590	14.07182	14.45998	13.96908	14.56272

```
plot(sarima.temps.2015, xlab = "Time", ylab = "Temperature (°C)", main =  
"Global Temperature ARIMA Forecast")  
lines(temps.ts)
```




```
plot(temps.ts)  
lines(sarima.temps.2015$mean, col = "green")
```




Analyze Residuals

```
plot(sarima.temps.mod$residuals)
```




```
acf(sarima.temps.mod$residuals)
```



Create function to calculate symmetric mean absolute percentage error (sMAPE) for forecast evaluation

```
sMAPE <- function(actual, estimate) {  
  absDev <- abs(estimate - actual)  
  return(sum(absDev/(estimate + actual))/length(actual))  
}  
  
(sMAPE.sarima.temps <- sMAPE(temps.test$temp, sarima.temps.2015$mean))  
## [1] 0.005514724
```

Holt-Winters

Time series with positive trend and seasonality

```
hw.temps.mod <- HoltWinters(temps.train.ts)  
summary(hw.temps.mod)
```

```
##           Length Class  Mode  
## fitted      2304   mts    numeric  
## x           588   ts     numeric  
## alpha        1  -none- numeric  
## beta         1  -none- numeric  
## gamma        1  -none- numeric  
## coefficients 14  -none- numeric
```


```
## seasonal      1  -none- character
## SSE           1  -none- numeric
## call          2  -none- call

hw.temps.mod

## Holt-Winters exponential smoothing with trend and additive seasonal
## component.
##
## Call:
## HoltWinters(x = temps.train.ts)
##
## Smoothing parameters:
##  alpha: 0.4205287
##  beta : 0
##  gamma: 0.2717648
##
## Coefficients:
##           [,1]
## a  15.983894264
## b   0.002918124
## s1 -1.803138843
## s2 -1.557478022
## s3 -0.817346858
## s4  0.055171372
## s5  0.791421348
## s6  1.363884431
## s7  1.621384127
## s8  1.569075627
## s9  0.955293615
## s10 0.045574377
## s11 -0.948840806
## s12 -1.632386848


plot(hw.temps.mod)
```

Holt-Winters filtering




```
plot(hw.temps.mod$fitted)
```

hw.temps.mod\$fitted



Analyzing Residuals

```
plot(hw.temps.mod$fitted[,1]-temps.train.ts)
```




white noise

Residuals seem like

```
acf(hw.temps.mod$fitted[,1]-temps.train.ts)
```

Series hw.temps.mod\$fitted[, 1] - temps.train.ts



No residual


autocorrelation

```
(hw.temps.2015 <- forecast(hw.temps.mod, h=12))
```

##	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2015		14.18367	14.04888	14.31846	13.97753	14.38982
## Feb 2015		14.43225	14.28603	14.57848	14.20862	14.65588
## Mar 2015		15.17530	15.01848	15.33213	14.93546	15.41515
## Apr 2015		16.05074	15.88398	16.21749	15.79571	16.30577
## May 2015		16.78991	16.61378	16.96603	16.52055	17.05927
## Jun 2015		17.36529	17.18027	17.55031	17.08232	17.64825
## Jul 2015		17.62571	17.43220	17.81922	17.32976	17.92165
## Aug 2015		17.57631	17.37467	17.77796	17.26793	17.88470
## Sep 2015		16.96545	16.75599	17.17491	16.64511	17.28579
## Oct 2015		16.05865	15.84166	16.27564	15.72679	16.39051
## Nov 2015		15.06715	14.84288	15.29142	14.72416	15.41015
## Dec 2015		14.38652	14.15520	14.61785	14.03274	14.74030

```
plot(hw.temps.2015, main = "Global Temperature Holt-Winters Forecast", xlab =
"Time", ylab = "Temperature (°C)")
lines(ts(temps$temp))
```

Global Temperature Holt-Winters Forecast



sMAPE Calculation

```
(sMAPE.hw.2015 <- sMAPE(temps.test$temp, hw.temps.2015$mean))
```

```
## [1] 0.003446822
```